

Complete Guide to Python Concurrency

Asyncio, Threading, and Multiprocessing

Comprehensive Learning Material

September 25, 2025

Contents

1	Introduction to Concurrency	2
1.1	Concurrency vs Parallelism	2
2	Asyncio - Cooperative Concurrency	2
2.1	Basic Concepts	2
2.2	Simple Asyncio Example	3
2.3	Concurrent Execution with <code>asyncio.gather()</code>	3
2.4	HTTP Requests with <code>aiohttp</code>	4
2.5	When to Use Asyncio	4
3	Threading - Preemptive Concurrency	5
3.1	Basic Threading Concepts	5
3.2	Thread Executors with Asyncio	5
3.3	Daemon vs Non-Daemon Threads	6
3.4	Race Conditions and Thread Safety	6
3.5	Deadlock Example	7
4	Multiprocessing - True Parallelism	8
4.1	Process Executors	8
4.2	Mixed Asyncio and Threading	9
5	Comparison and Decision Guide	9
6	Decision Flowchart	10
7	Best Practices	10
7.1	Asyncio Best Practices	10
7.2	Threading Best Practices	10
7.3	General Guidelines	10
8	Performance Analysis	11
8.1	Execution Time Comparison	11
8.2	Resource Usage	11
9	Conclusion	11

1 Introduction to Concurrency


Key Concept

Concurrency is the ability to handle multiple tasks seemingly at the same time. In Python, we have three main approaches:

- **Asyncio**: Cooperative concurrency (single-threaded)
- **Threading**: Preemptive concurrency (multi-threaded)
- **Multiprocessing**: True parallelism (multi-process)

1.1 Concurrency vs Parallelism

Concurrency (Time-slicing)



Parallelism (Simultaneous)



2 Asyncio - Cooperative Concurrency

2.1 Basic Concepts

Key Concept

Asyncio uses an event loop to manage coroutines. Key concepts:

- **Coroutines**: Functions defined with `async def`
- **Await**: Pauses execution and yields control
- **Event Loop**: Manages and executes coroutines
- **Tasks**: Scheduled coroutines

2.2 Simple Asyncio Example

Code Example

```
1 import asyncio
2
3 async def brew_chai():
4     print("Brewing chai...")
5     await asyncio.sleep(2) # Non-blocking sleep
6     print("Chai is ready")
7
8 # Run the coroutine
9 asyncio.run(brew_chai())
```

What this code does:


1. Defines an async function (coroutine)
2. Uses `await asyncio.sleep(2)` instead of blocking `time.sleep()`
3. `asyncio.run()` creates event loop and runs the coroutine
4. During the 2-second wait, control can be yielded to other tasks

2.3 Concurrent Execution with `asyncio.gather()`

Code Example

```
1 import asyncio
2
3 async def brew(name):
4     print(f"Brewing {name}...")
5     await asyncio.sleep(3)
6     print(f"{name} is ready...")
7
8 async def main():
9     # Run multiple coroutines concurrently
10    await asyncio.gather(
11        brew("Masala chai"),
12        brew("Green chai"),
13        brew("Ginger chai"),
14    )
15
16 asyncio.run(main())
```

Execution Flow:



2.4 HTTP Requests with aiohttp

Code Example

```
1 import asyncio
2 import aiohttp
3
4 async def fetch_url(session, url):
5     async with session.get(url) as response:
6         print(f"Fetches {url} with status {response.status}")
7
8 async def main():
9     urls = ["https://httpbin.org/delay/2"] * 3
10    async with aiohttp.ClientSession() as session:
11        tasks = [fetch_url(session, url) for url in urls]
12        await asyncio.gather(*tasks)
13
14 asyncio.run(main())
```

Key Points:

- aiohttp provides async HTTP client
- `async with` ensures proper resource cleanup
- All 3 requests run concurrently, completing in 2 seconds total
- `*tasks` unpacks the list for `gather()`

2.5 When to Use Asyncio

Best For	Advantages	Limitations
I/O bound tasks Network re-requests File operations	Memory efficient No GIL issues	Single-threaded Requires async libraries
Database queries	Easy to reason about Built-in event loop	CPU-bound tasks blocked Learning curve

Table 1: Asyncio Use Cases

3 Threading - Preemptive Concurrency

3.1 Basic Threading Concepts

Key Concept

Threading allows multiple threads to run within a single process:

- **Thread**: Lightweight execution unit
- **GIL**: Global Interpreter Lock (limits true parallelism)
- **Daemon threads**: Die when main program exits
- **Race conditions**: Concurrent access to shared data

3.2 Thread Executors with Asyncio

Code Example

```
1 import asyncio
2 import time
3 from concurrent.futures import ThreadPoolExecutor
4
5 def check_stock(item):
6     print(f"Checking {item} in store...")
7     time.sleep(3) # Blocking operation
8     return f"{item} stock: 42"
9
10 async def main():
11     loop = asyncio.get_running_loop()
12     with ThreadPoolExecutor() as pool:
13         # Run blocking function in thread pool
14         result = await loop.run_in_executor(pool, check_stock, "
15             Masala chai")
16         print(result)
17 asyncio.run(main())
```

What happens here:

1. `check_stock()` is a blocking function (uses `time.sleep()`)
2. `ThreadPoolExecutor` manages a pool of worker threads
3. `run_in_executor()` runs the function in a separate thread
4. Main `asyncio` event loop remains responsive

3.3 Daemon vs Non-Daemon Threads

Code Example

```

1 import threading
2 import time
3
4 def monitor_tea_temp():
5     while True:
6         print("Monitoring tea temperature...")
7         time.sleep(2)
8
9 # Daemon thread - dies with main program
10 t_daemon = threading.Thread(target=monitor_tea_temp, daemon=True)
11 t_daemon.start()
12
13 # Non-daemon thread - keeps program alive
14 t_normal = threading.Thread(target=monitor_tea_temp)
15 t_normal.start()
16
17 print("Main program done")

```

Thread Type	Daemon=True	Daemon=False
Behavior	Dies when main exits	Keeps program running
Use case	Background monitoring	Critical operations
Example	Logging, heartbeats	Data processing

Table 2: Daemon vs Non-Daemon Threads

3.4 Race Conditions and Thread Safety

Code Example

```

1 import threading
2
3 chai_stock = 0
4
5 def restock():
6     global chai_stock
7     for _ in range(100000):
8         chai_stock += 1 # This is NOT atomic!
9
10 threads = [threading.Thread(target=restock) for _ in range(2)]
11
12 for t in threads:
13     t.start()
14 for t in threads:
15     t.join()
16
17 print("Chai stock:", chai_stock) # Often < 200000!

```


Why the result is wrong:

1. `chai_stock += 1` involves: read \rightarrow increment \rightarrow write
2. Two threads can read the same value simultaneously
3. Both increment and write back, losing one increment
4. This is a **race condition**

3.5 Deadlock Example

Code Example

```
1 import threading
2
3 lock_a = threading.Lock()
4 lock_b = threading.Lock()
5
6 def task1():
7     with lock_a:
8         print("Task 1 acquired lock a")
9         with lock_b: # Waits for lock_b
10             print("Task 1 acquired lock b")
11
12 def task2():
13     with lock_b:
14         print("Task 2 acquired lock b")
15         with lock_a: # Waits for lock_a
16             print("Task 2 acquired lock a")
17
18 t1 = threading.Thread(target=task1)
19 t2 = threading.Thread(target=task2)
20
21 t1.start()
22 t2.start() # This will deadlock!
```



DEADLOCK!
Each waits for other's lock

4 Multiprocessing - True Parallelism

4.1 Process Executors

Code Example

```
1 import asyncio
2 from concurrent.futures import ProcessPoolExecutor
3
4 def encrypt(data):
5     # CPU-intensive operation
6     return f"    {data[::-1]}"
7
8 async def main():
9     loop = asyncio.get_running_loop()
10    with ProcessPoolExecutor() as pool:
11        result = await loop.run_in_executor(pool, encrypt, "
12            credit_card_1234")
13        print(result)
14
15 if __name__ == "__main__":
16     asyncio.run(main())
```

Key differences from threading:

- Separate Python interpreter per process
- No GIL limitations
- Higher memory overhead
- Inter-process communication needed for shared data

4.2 Mixed Asyncio and Threading

Code Example

```
1 import asyncio
2 import threading
3 import time
4
5 def background_worker():
6     while True:
7         time.sleep(1)
8         print("Logging system health")
9
10 async def fetch_orders():
11     await asyncio.sleep(3)
12     print("Order fetched")
13
14 # Start background thread
15 threading.Thread(target=background_worker, daemon=True).start()
16
17 # Run async main function
18 asyncio.run(fetch_orders())
```

This pattern combines:

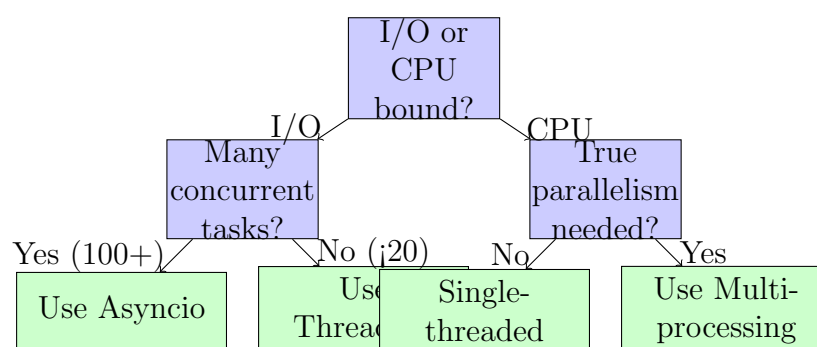
1. Background monitoring thread (daemon)
2. Asyncio event loop for I/O operations
3. Non-blocking cooperation between both

5 Comparison and Decision Guide

Aspect	Asyncio	Threading	Multiprocessing
Best for	I/O bound tasks	Mixed I/O + some CPU	CPU bound tasks
Memory	Very efficient	Moderate overhead	High overhead
GIL Impact	No impact	Limited by GIL	No GIL
Communication	Shared memory	Shared memory + locks	IPC mechanisms
Debugging	Moderate difficulty	Hard (race conditions)	Harder (separate processes)
Scalability	Thousands of tasks	Dozens of threads	Number of CPU cores
Libraries	Requires async libs	Standard libraries work	Standard libraries work

Table 3: Concurrency Method Comparison

6 Decision Flowchart



7 Best Practices

7.1 Asyncio Best Practices

- Always use `async with` for resource management
- Don't mix blocking and async code
- Use `asyncio.gather()` for concurrent tasks
- Handle exceptions in coroutines properly
- Use connection pooling for HTTP clients

7.2 Threading Best Practices

- Use `ThreadPoolExecutor` instead of raw threads
- Always use locks for shared mutable data
- Prefer daemon threads for background tasks
- Avoid circular lock dependencies
- Use `queue.Queue` for thread communication

7.3 General Guidelines

- Start simple - single-threaded first
- Profile before optimizing
- Consider the GIL for CPU-bound tasks
- Use appropriate data structures (thread-safe)
- Plan error handling and resource cleanup

8 Performance Analysis

8.1 Execution Time Comparison

Task Type	Sequential	Asyncio	Threading	Multiprocessing
3 HTTP requests (2s each)	6s	2s	2s	2s
File I/O operations	10s	3s	4s	3s
CPU-intensive calculation	8s	8s	8s*	2s
Mixed I/O + CPU	15s	10s	12s	6s

Table 4: Performance Comparison (*GIL limited)

8.2 Resource Usage

9 Conclusion

Understanding Python's concurrency models is crucial for building efficient applications:

- **Asyncio:** Perfect for I/O-heavy applications with many concurrent operations
- **Threading:** Good for mixed workloads with occasional blocking operations
- **Multiprocessing:** Essential for CPU-bound tasks requiring true parallelism

The key is matching the right tool to your specific use case. Start with profiling your application to understand whether you're I/O or CPU bound, then choose the appropriate concurrency model.

Remember: "Premature optimization is the root of all evil" - start simple and optimize when you have real performance requirements and measurements.